

# HDFS Permissions

Security is an important consideration for any computer system, and big data systems are no exception. In the VM for this course, you have all the permissions you need to do the exercises and examples in this course, but at your company, you may not be able to do everything you're learning here. For example, the ability to create databases, or even tables, is often restricted so a company's system doesn't get overloaded with redundant or otherwise unnecessary items. Similarly, the ability to drop tables, or to delete or move files, may be restricted. Larger companies will only allow their data engineers, database administrators, or system administrators to do this work.

A deep look at permission issues is beyond the scope of this course, but the following resources can provide additional information, if you're interested:

1. DWGeek's Hadoop Security – [Hadoop HDFS File Permissions](#) is a quick overview of file permissions. This is a good place to start, and then you can decide if you want to dig deeper.
2. [HDFS Commands, HDFS Permissions and HDFS Storage](#) is a chapter from a book, *Expert Hadoop Administration: Managing, Tuning, and Securing Spark, YARN, and HDFS*. The section “HDFS Users and Superusers” may be helpful if you are unfamiliar with the concept of a superuser.
3. The documentation on Hadoop includes the [HDFS Permissions Guide](#), which likely has more detail than you'll ever want, but it's a great resource if there's a particular question you're looking to answer.