

# Using Sqoop to Export Data

You not only can import data from an RDBMS into Hadoop, but you can send data the other way as well, using the Sqoop export command.

For example, suppose some new product recommendations have been generated after some processing on the Hadoop cluster. These recommendations need to be exported to the web site's back-end database. This can be done with the following command:

```
$ sqoop export \  
  --connect jdbc:mysql://localhost/company \  
  --username jdoe \  
  --password bigsecret \  
  --table product_recommendations \  
  --export-dir /mydata/recommender_output
```

The `--export-dir` argument specifies where the data to be exported is located in HDFS; in this case, it's in the `/mydata/recommender_output` directory. The destination table in the RDBMS is identified in the `--table` option; in this case, the table will be `product_recommendations`. Note that this only *exports the data*, it doesn't *create the table* in the RDBMS—the destination table must already exist.

While Sqoop will let you import all tables into HDFS using a single command, it does not have a command to export more than one table. Exporting must be done one table directory at a time.

For more options, see the documentation through <https://sqoop.apache.org>. Be sure to use the documentation for the version you are using. To see what version of Sqoop you are using, run the command `sqoop version`.